

# Mercedes-Benz Greener Manufacturing

The Personal Case Study

## 1. Business Problem



### 1.1 Description

Since the first automobile, the Benz Patent Motor Car in 1886, Mercedes-Benz has stood for important automotive innovations. These include, for example, the passenger safety cell with crumple zone, the airbag and intelligent assistance systems. Mercedes-Benz applies for nearly 2000 patents per year, making the brand the European leader among premium car makers. Daimler's Mercedes-Benz cars are leaders in the premium car industry. With a huge selection of features and options, customers can choose the customized Mercedes-Benz of their dreams.

To ensure the safety and reliability of each and every unique car configuration before they hit the road, Daimler's engineers have developed a robust testing system. But, optimizing the speed of their testing system for so many possible feature combinations is complex and time-consuming without a powerful algorithmic approach. As one of the world's biggest manufacturers of premium cars, safety and efficiency are paramount on Daimler's production lines.

### 1.2 Data Description:

This dataset contains an anonymized set of variables, each representing a custom feature in a Mercedes car. For example, a variable could be 4WD, added air suspension, or a head-up display.

The ground truth is labeled 'y' and represents the time (in seconds) that the car took to pass testing for each variable.

## 1.3 Objective

1. Reduce the time that cars spend on the test bench.
2. To speedier testing, resulting in lower carbon dioxide emissions without reducing Daimler's standards.

## 1.4 Sources/Useful Links

Main Source: <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/overview/description>  
(<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/overview/description>).

Data: <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data> (<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data>).

Discussion: <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/discussion>  
(<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/discussion>).

## 2. Machine Learning Problem

### 2.1 Data

#### 2.1.1 Data Overview

Get the data from : <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data>  
(<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data>)

- Data will be in a file Train.csv and Test.csv
- Size of Train.csv - 3.07MB
- Size of Train.csv - 3.04 MB
- Number of rows and columns in Train.csv = 4210 x 378
- Number of rows and columns in Test.csv = 4210 x 377

#### 2.1.2 Example Data point

[illegible]

- `sklearn.metrics.r2_score(y_true, y_pred, sample_weight=None, multioutput='uniform_average')`
- $R^2$  (coefficient of determination) regression score function.
- Best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of  $y$ , disregarding the input features, would get a  $R^2$  score of 0.0.

## 2.3 Train and Test Construction

We build train and test by randomly splitting in the ratio of 70:30 or 80:20 whatever we choose as we have sufficient points to work with.

# 3. Exploratory Data Analysis

## 3.1 Importing important libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
import xgboost as xgb
color = sns.color_palette()
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA
from sklearn.metrics import r2_score
%matplotlib inline
from sklearn.svm import SVR, LinearSVC
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings("ignore")
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
import scipy.stats as stats
from sklearn.externals import joblib
```

## 3.1 Reading data and basic stats

```
In [2]: # Load the Drive helper and mount
#from google.colab import drive
# This will prompt for authorization.
#drive.mount('/content/drive')
```

```
In [3]: #Loading data from google drive
#train_df = pd.read_csv('/content/drive/My Drive/Colab Notebooks/Personal Case Study/train.csv')
#print("Train shape : ", train_df.shape)
#test_df = pd.read_csv('/content/drive/My Drive/Colab Notebooks/Personal Case Study/test.csv')
#print("Train shape : ", test_df.shape)
```

```
In [4]: #Loading data from HDD
train_df = pd.read_csv('train.csv')
print("Train shape : ", train_df.shape)
test_df = pd.read_csv('test.csv')
print("Train shape : ", test_df.shape)
```

```
Train shape : (4209, 378)
Train shape : (4209, 377)
```

```
In [5]: train_df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB
```

```
In [6]: train_df.head()
```

```
Out[6]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382	X385
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0	0

5 rows × 378 columns



### Target Variable:

- "y" is the variable we need to predict. So let us do some analysis on this variable first.
- Variable y is of type float
- X0,X1,X2,X3,X4,X5,X6,X8 are of type object
- Rest of the columns are int type
- We will convert [X0,X1,X2,X3,X4,X5,X6,X8] to categorical types and plot to see the distribution of values.

## 3.2 Checking for missing values

```
In [6]: def check_missing_values(df):

        if df.isnull().any().any():
            print("There are missing values in the data")
        else:
            print("There are no missing values in the data")
```

```
In [7]: #calling functions to check missing values on training and test datasets
check_missing_values(train_df)
check_missing_values(test_df)
```


```
There are no missing values in the data
There are no missing values in the data
```

## 3.3 Plotting

### 3.3.1 Plotting y values

```
In [8]: #we are checking 'y' column
plt.figure(figsize=(8,6))
plt.scatter(range(train_df.shape[0]), np.sort(train_df.y.values))
plt.xlabel('index', fontsize=12)
plt.ylabel('y', fontsize=12)
plt.show()
```


""""here we have observed 1 outlier at apporx 260""""



Out[8]: 'here we have observed 1 outlier at apporx 260'

```
In [9]: # we again check by visualising in BoxPlot
```

```
plt.figure(figsize=(15,5))
sns.boxplot(train_df.loc[:, 'y'])
plt.show()
```



here we have observed 1 outlier


```
In [10]: # we need to remove that outlier
# i have used zscore method and set threshold 10 acc to our data
# https://www.geeksforgeeks.org/scipy-stats-zscore-function-python/

train_df['x'] = np.abs(stats.zscore(train_df.loc[:, 'y']))

outlier_ids = train_df[train_df['x'] > 10].ID


train_df_final = train_df[~train_df['ID'].isin(list(outlier_ids))]
```

```
In [11]: #now plotting again without outlier
#we are checking 'y' column
plt.figure(figsize=(8,6))
plt.scatter(range(train_df_final.shape[0]), np.sort(train_df_final.y.values))
plt.xlabel('index', fontsize=12)
plt.ylabel('y', fontsize=12)
plt.show()
```



In [12]: *# we again check by visualising in BoxPlot*


```
plt.figure(figsize=(15,5))
sns.boxplot(train_df_final.loc[:, 'y'])
plt.show()
```



### 3.3.2 Plotting y distribution graph.

In [13]: *ulimit = 180# we have taken 180 data points*  
*train\_df\_final['y'].ix[train\_df\_final['y']>ulimit] = ulimit*

```
plt.figure(figsize=(12,8))#plot size
sns.distplot(train_df_final.y.values, bins=50, kde=True)
plt.xlabel('y value', fontsize=12)
plt.show()
```



In [14]: *#removing that x helper row for outlier from main row*

```
train_df_final = train_df_final.drop(["x"], axis=1)
```



### 3.3.3 Data type of all the variables present in the dataset.

```
In [15]: dtype_df = train_df_final.dtypes.reset_index()
dtype_df.columns = ["Count", "Column Type"]
dtype_df.groupby("Column Type").aggregate('count').reset_index()
```

Out[15]:

	Column Type	Count
0	int64	369
1	float64	1
2	object	8

Maximum of the columns are integers.

8 categorical columns.

1 float column (target variable) i.e. 'y'

```
In [16]: #here we can see their types
dtype_df.ix[:15,:]
```

Out[16]:


	Count	Column Type
0	ID	int64
1	y	float64
2	X0	object
3	X1	object
4	X2	object
5	X3	object
6	X4	object
7	X5	object
8	X6	object
9	X8	object
10	X10	int64
11	X11	int64
12	X12	int64
13	X13	int64
14	X14	int64
15	X15	int64

X0 to X8 are the categorical columns


### 3.4 Plotting these categorical Values

```
In [17]: #https://www.kaggle.com/sudalairajkumar/simple-exploration-notebook-mercedes
var_name = ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']
for val in var_name:
    col_order = np.sort(train_df_final[val].unique()).tolist()
    plt.figure(figsize=(12,6))
    sns.stripplot(x=val, y='y', data=train_df_final, order=col_order)
    plt.xlabel(val, fontsize=12)
    plt.ylabel('y', fontsize=12)
    plt.title("Distribution of y variable with "+val, fontsize=15)
    plt.show()
```


Distribution of y variable with X0




Distribution of y variable with X1




Distribution of y variable with X2




Distribution of y variable with X3





Distribution of y variable with X4



Distribution of y variable with X5





Observation:

- we have observed that X0, X1, X2, X5, X6 and X8 have larger data point.
- X4 and X3 have lesser data point.

## 4. Machine Learning Models

### 4.1 Data preparation

1 one of doing is this

```
In [2]: train_df = pd.read_csv('train.csv')
print("Train shape : ", train_df.shape)
test_df = pd.read_csv('test.csv')
print("Train shape : ", test_df.shape)

y_train = train_df['y'].values
id_test = test_df['ID'].values

usable_columns = list(set(train_df.columns) - set(['ID', 'y']))#taking only important columns
print(len(usable_columns))

x_train_final = train_df[usable_columns]
x_test_final = test_df[usable_columns]
```

Train shape : (4209, 378)  
Train shape : (4209, 377)  
376

```
In [3]: # Converting training dataset object categorical values to numerical categorical type
s
#taken help from Link: https://www.kaggle.com/anokas/mercedes-edxgboost-starter-0-5
5

for column in usable_columns:
    cardinality = len(np.unique(x_train_final[column]))

    if cardinality == 1:
        x_train_final.drop(column, axis=1) # Column with only one value is useless so
we drop it.
        x_test_final.drop(column, axis=1)

    if cardinality > 2: # Column is categorical.
        mapper = lambda x: sum([ord(digit) for digit in x])
        x_train_final[column] = x_train_final[column].apply(mapper)
        x_test_final[column] = x_test_final[column].apply(mapper)
```

```
In [4]: # spilting it into 70:30 ratio
X_train, X_test, y_train, y_test = train_test_split(x_train_final, y_train, test_size=0.3, random_state=42)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(2946, 376)  
(1263, 376)  
(2946,)  
(1263,)

## XGBoost

```
In [5]: #taken help from kaggle discussion and kernels for xgboost
#setting up xtrain and xtrain
```

```
y_mean = y_train.mean()

d_train = xgb.DMatrix(X_train, label=y_train)
d_cvalid = xgb.DMatrix(X_test, label=y_test)
d_test = xgb.DMatrix(x_test_final)
```

```
In [6]: # evaluation r2_score metric
def r2_score_metric(y_pred, y):
    y_true = y.get_label()
    return 'r2', r2_score(y_true, y_pred)
```

```
In [159]: %%time
#xgb parameters
#just cross validation our model

params = {

    'eta': 0.005,
    'learning_rate': 0.008,
    'max_depth': 4,
    'subsample': 0.9,
    'objective': 'reg:linear',
    'n_estimators': 687,
    'eval_metric': 'rmse',
    'base_score': y_mean, # base prediction = mean(target)
    'silent': 1
}

num_boost_round=2000

#Cross Validation of XGBoost
cv_result = xgb.cv(params,
                    d_train,
                    num_boost_round,
                    nfold = 3,
                    early_stopping_rounds=50,
                    feval=r2_score_metric,#here we have used our metric method
                    verbose_eval=100,
                    show_stdv=False
                    )

[0]      train-r2:0.00950233      train-rmse:12.3021      test-r2:0.00846033      test
-rmse:12.3007
[100]    train-r2:0.502523      train-rmse:8.71869      test-r2:0.478775      test
-rmse:8.92026
[200]    train-r2:0.612515      train-rmse:7.69454      test-r2:0.571492      test
-rmse:8.08772
[300]    train-r2:0.648009      train-rmse:7.3335      test-r2:0.587773      test
-rmse:7.93169
[400]    train-r2:0.666977      train-rmse:7.133      test-r2:0.590321      test
-rmse:7.90657
Wall time: 34.2 s
```

**best cv\_result is: train-r2:0.67011 test-r2:0.591521**

```
In [161]: %%time
#Training the model
#taken help from link: https://www.kaggle.com/anokas/mercedes-edxgboost-starter-0-55

#model = joblib.load('model_xgb.pkl')#from load

watchlist = [(d_train, 'train'), (d_cvalid, 'valid')]

model = xgb.train(params, d_train , num_boost_round, watchlist, early_stopping_rounds=60,
                  feval=r2_score_metric, maximize=True, verbose_eval=10)

#joblib.dump(model, 'model_xgb.pkl')#to load
```



[0] train-rmse:12.3026 valid-rmse:13.333 train-r2:0.009819 vali  
d-r2:0.005644  
Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 60 rounds.

[10]	train-rmse:11.7348	valid-rmse:12.8415	train-r2:0.099115	vali
d-r2:0.077599				
[20]	train-rmse:11.2266	valid-rmse:12.4077	train-r2:0.17545	vali
d-r2:0.138868				
[30]	train-rmse:10.7723	valid-rmse:12.0268	train-r2:0.240831	vali
d-r2:0.190928				
[40]	train-rmse:10.3694	valid-rmse:11.6967	train-r2:0.296558	vali
d-r2:0.234734				
[50]	train-rmse:10.0114	valid-rmse:11.4052	train-r2:0.344297	vali
d-r2:0.272398				
[60]	train-rmse:9.69455	valid-rmse:11.1513	train-r2:0.385139	vali
d-r2:0.304428				
[70]	train-rmse:9.41513	valid-rmse:10.9318	train-r2:0.420072	vali
d-r2:0.33154				
[80]	train-rmse:9.16923	valid-rmse:10.7439	train-r2:0.449969	vali
d-r2:0.354329				
[90]	train-rmse:8.95299	valid-rmse:10.582	train-r2:0.475606	vali
d-r2:0.373644				
[100]	train-rmse:8.76199	valid-rmse:10.4408	train-r2:0.497742	vali
d-r2:0.390249				
[110]	train-rmse:8.59695	valid-rmse:10.3209	train-r2:0.516485	vali
d-r2:0.404165				
[120]	train-rmse:8.45156	valid-rmse:10.2174	train-r2:0.532701	vali
d-r2:0.416065				
[130]	train-rmse:8.32342	valid-rmse:10.1303	train-r2:0.546763	vali
d-r2:0.425975				
[140]	train-rmse:8.21268	valid-rmse:10.0551	train-r2:0.558744	vali
d-r2:0.434468				
[150]	train-rmse:8.11545	valid-rmse:9.99483	train-r2:0.56913	vali
d-r2:0.441223				
[160]	train-rmse:8.03094	valid-rmse:9.94416	train-r2:0.578057	vali
d-r2:0.446874				
[170]	train-rmse:7.95588	valid-rmse:9.90191	train-r2:0.585907	vali
d-r2:0.451564				
[180]	train-rmse:7.88971	valid-rmse:9.86327	train-r2:0.592766	vali
d-r2:0.455835				
[190]	train-rmse:7.83187	valid-rmse:9.83162	train-r2:0.598716	vali
d-r2:0.459322				
[200]	train-rmse:7.77966	valid-rmse:9.80499	train-r2:0.604048	vali
d-r2:0.462248				
[210]	train-rmse:7.73178	valid-rmse:9.78603	train-r2:0.608907	vali
d-r2:0.464325				
[220]	train-rmse:7.68943	valid-rmse:9.7693	train-r2:0.61318	vali
d-r2:0.466155				
[230]	train-rmse:7.65265	valid-rmse:9.75688	train-r2:0.616871	vali
d-r2:0.467512				
[240]	train-rmse:7.6196	valid-rmse:9.74591	train-r2:0.620173	vali
d-r2:0.468708				
[250]	train-rmse:7.58943	valid-rmse:9.73683	train-r2:0.623175	vali
d-r2:0.469698				
[260]	train-rmse:7.56162	valid-rmse:9.72916	train-r2:0.625931	vali
d-r2:0.470533				
[270]	train-rmse:7.5358	valid-rmse:9.72583	train-r2:0.628481	vali
d-r2:0.470896				
[280]	train-rmse:7.51009	valid-rmse:9.72262	train-r2:0.631013	vali
d-r2:0.471245				
[290]	train-rmse:7.48662	valid-rmse:9.71946	train-r2:0.633315	vali
d-r2:0.471589				
[300]	train-rmse:7.46653	valid-rmse:9.71587	train-r2:0.63528	vali
d-r2:0.471979				

[310]	train-rmse:7.44646	valid-rmse:9.71569	train-r2:0.637238	vali
	d-r2:0.471998			
[320]	train-rmse:7.42855	valid-rmse:9.71446	train-r2:0.638981	vali
	d-r2:0.472132			
[330]	train-rmse:7.41077	valid-rmse:9.71458	train-r2:0.640708	vali
	d-r2:0.472119			
[340]	train-rmse:7.39414	valid-rmse:9.71491	train-r2:0.642319	vali
	d-r2:0.472083			
[350]	train-rmse:7.38027	valid-rmse:9.71499	train-r2:0.643659	vali
	d-r2:0.472075			
[360]	train-rmse:7.3651	valid-rmse:9.71774	train-r2:0.645122	vali
	d-r2:0.471775			
[370]	train-rmse:7.35246	valid-rmse:9.71885	train-r2:0.646339	vali
	d-r2:0.471654			
[380]	train-rmse:7.33871	valid-rmse:9.71988	train-r2:0.647661	vali
	d-r2:0.471542			
Stopping. Best iteration:				
[327]	train-rmse:7.41692	valid-rmse:9.71405	train-r2:0.640111	vali
	d-r2:0.472177			

Wall time: 15.7 s

In [162]: *# Predict on test*

```
y_pred = model.predict(d_test)
```


In [ ]: *# Predicting R2SCORE*

```
#from sklearn.metrics import r2_score
#r2_score = r2_score(y_test, y_pred)#taking r2score on traing data
#print('r2_score = ',r2_score)
```

In [163]: *#exporting final results into csv file*

```
csvfile = pd.DataFrame()
csvfile['ID'] = test_df['ID']
csvfile['y'] = y_pred
csvfile.to_csv('xgbcsv11.csv', index=False)
```

**Final Test result given by kaggle is: 0.54981**

 Featured Prediction Competition

## Mercedes-Benz Greener Manufacturing

Can you cut the time a Mercedes-Benz spends on the test bench?

**\$25,000**

Prize Money



Daimler · 3,835 teams · 2 years ago

[Overview](#)

[Data](#)

[Kernels](#)

[Discussion](#)

[Leaderboard](#)

[Rules](#)

[Team](#)

[My Submissions](#)

[Late Submission](#)

### Your most recent submission

Name	Submitted	Wait time	Execution time	Score
xgbcsv11.csv	just now	0 seconds	0 seconds	0.54981


Complete

[Jump to your position on the leaderboard](#) ▼

## Feature Importance

```
In [76]: #https://www.kaggle.com/satadru5/mercedes-benz-xgb-modeling-lb-score-0-54472
fig, ax = plt.subplots(1, 1, figsize=(8, 13))
xgb.plot_importance(model, max_num_features=50, height=0.5, ax=ax)
```


Out[76]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f5e09445d30>



## 4.2 PCA - Principal component analysis

```
In [0]: # PCA Implementation
pca = PCA(n_components=2)
pca_data = pca.fit_transform(X_train)
```

```
In [0]: cmap = sns.cubehelix_palette(as_cmap=True)
f, ax = plt.subplots(figsize=(10,7))
points = ax.scatter(pca_data[:,0], pca_data[:,1], c=y_train, s=50, cmap=cmap)
f.colorbar(points)
plt.show()
```




Observation: here we can see how PCA visualize all the data point far separated from each other, they are not forming tightly group.

### 4.3 T-SNE (t-distributed Stochastic Neighbor Embedding)

```
In [0]: # TSNE Implementation
model = TSNE(n_components=2,random_state=0,perplexity=30)

tsne_data = model.fit_transform(X_train)
```

```
In [0]: cmap = sns.cubehelix_palette(as_cmap=True)
f, ax = plt.subplots(figsize=(10,7))
points = ax.scatter(tsne_data[:,0], tsne_data[:,1], c=y_train, s=50, cmap=cmap)
f.colorbar(points)
plt.show()
```



Observation: here we can see how TSNE visualize all the data point closely attached from each other, They are well grouped .

## 4.4 K-Nearest Neighbors Regressor

```
In [0]: #KNN implementation
#biulding model

knn = KNeighborsRegressor(n_neighbors=5)#k=5 gives best results

knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

r2_score_knn = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(knn.score(X_train, y_train) *100,2)#taking accuracy

results = {'r2_score':r2_score_knn, 'accuracy':accuracy}
print (results)

{'r2_score': 0.414, 'accuracy': 65.57}
```

## 4.5 Support Vector Regressor

```
In [0]: #SVR implementation
from sklearn.metrics import r2_score
clf = SVR()

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

r2_score = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(clf.score(X_train, y_train) * 100, 2)

results = {'r2_score':r2_score, 'accuracy':accuracy}
print(results)

{'r2_score': 0.384, 'accuracy': 44.76}
```

## 4.6 Random Forest Regressor

```
In [0]: #RFR implementation
from sklearn.metrics import r2_score
clf = RandomForestRegressor(n_estimators = 60 ,max_depth=5,oob_score=True)

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

r2_score = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(clf.score(X_train, y_train) * 100, 2)

results = {'r2_score':r2_score, 'accuracy':accuracy}
print (results)

{'r2_score': 0.533, 'accuracy': 65.66}
```

## 4.7 Linear Regression

```
In [143]: #Linear Regression implementation
from sklearn.metrics import r2_score
clf = LinearRegression()

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

r2_score = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(clf.score(X_train, y_train) * 100, 2)

results = {'r2_score':r2_score, 'accuracy':accuracy}
print (results)

{'r2_score': 0.439, 'accuracy': 63.06}
```

## 5. Conclusion

S.no	Model Algo	R2 Score	Accuracy
1.	K-Nearest Neighbors Regressor	0.414	65.57
2.	Support Vector Regressor	0.384	44.76
3.	Random Forest Regressor	0.533	65.66
3.	Linear Regression	0.439	63.06
4.	XGBoost	0.54981	-

Here we can see from conclusion that

***1. XGBoost perform best in this case by obtaining R2\_SCORE =0.54981***

***2.Next Our RANDOM FOREST REGRESSOR also perform well with nearly R2\_SCORE=0.533***

***3. After that Knn does better.***

***4. Linear Regression model not suit for this type of problem.***